


The Modular Framework for Scalable Multi-Agent Orchestration



The Issue

With the surge of Agentic AI, **traditional systems** struggle to keep up



Our Solution

A system that coordinates multiple AI agents to **work together automatically**, to dynamically tackle complex problems




Target Audience

Enterprise organizations **seeking scalable AI systems** that can automate complex customer service and operational workflows.


Tech Stack




System Architecture & How it Works



Universal Delivery One Backend, Multiple Interfaces.



The Orchestrator can seamlessly serve data to any client interface
(Chatbot Interfaces, SMS/Email Notifications, Proactive Insight Dashboards)



Observability as a Feature

Deep-dive observability logs every LLM reasoning step, ensuring total transparency and compliance for enterprise deployment.

Acknowledgments

We would like to sincerely thank **Professor Ian Harris** and the **Accenture** team, especially **Shawna Tuli**, **Vishrut Chokshi**, and **Mo Nomeli**, for their continued support and guidance.

Add Us on LinkedIn!



Harish S.



Arya G.



Sanket L.



Shaun L.

Results vs. Baselines

(Percentages compiled on 35 Test Cases over 210 runs)

Metric	One-shot LLM	RAG+LLM	Our System
Task Success	82%	85%	93%
Fact Coverage	72%	80%	92%
Sequence Accuracy	N/A	N/A	96%
Avg Latency	3.0s	3.7s	*17.1s

*~66% of this is LLM planning + synthesis; the remaining ~34% is sequential agent execution.

Key Takeaways

- Orchestration **improved** multi-hop reasoning
- Dynamic agents **outperformed** static retrieval
- Registry-based design enabled **plug-and-play** extensibility

Why Orchestration Matters

Single-call systems handle simple retrieval but **break on multi-step workflows**.

Our orchestrator's **largest wins** are on **anomaly detection** (+25% fact coverage), **billing** (+15%), and **weather** (+15%) — the categories that require composing multiple sources.